


```
In [5]: import numpy as np
import pandas as pd
%matplotlib inline
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [20]: dataset = pd.read_csv("HR_Employee_Attrition_IBM.csv")
dataset.head()
```

Out[20]:

	Age	Attrition	BusinessTravel	DailyRate	Department	DistanceFromHome	Education	EducationField	EmployeeCount	EmployeeN
0	41	Yes	Travel_Rarely	1102	Sales	1	2	Life Sciences	1	
1	49	No	Travel_Frequently	279	Research & Development	8	1	Life Sciences	1	
2	37	Yes	Travel_Rarely	1373	Research & Development	2	2	Other	1	
3	33	No	Travel_Frequently	1392	Research & Development	3	4	Life Sciences	1	
4	27	No	Travel_Rarely	591	Research & Development	2	1	Medical	1	

5 rows × 35 columns



In [9]: `dataset.info()`

```


<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1470 entries, 0 to 1469
Data columns (total 35 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Age                                   1470 non-null   int64
1   Attrition                           1470 non-null   object
2   BusinessTravel                      1470 non-null   object
3   DailyRate                           1470 non-null   int64
4   Department                          1470 non-null   object
5   DistanceFromHome                    1470 non-null   int64
6   Education                           1470 non-null   int64
7   EducationField                      1470 non-null   object
8   EmployeeCount                       1470 non-null   int64
9   EmployeeNumber                      1470 non-null   int64
10  EnvironmentSatisfaction              1470 non-null   int64
11  Gender                              1470 non-null   object
12  HourlyRate                          1470 non-null   int64
13  JobInvolvement                      1470 non-null   int64
14  JobLevel                            1470 non-null   int64
15  JobRole                             1470 non-null   object
16  JobSatisfaction                     1470 non-null   int64
17  MaritalStatus                       1470 non-null   object
18  MonthlyIncome                       1470 non-null   int64
19  MonthlyRate                         1470 non-null   int64
20  NumCompaniesWorked                  1470 non-null   int64
21  Over18                              1470 non-null   object
22  OverTime                            1470 non-null   object
23  PercentSalaryHike                   1470 non-null   int64
24  PerformanceRating                   1470 non-null   int64
25  RelationshipSatisfaction             1470 non-null   int64
26  StandardHours                       1470 non-null   int64
27  StockOptionLevel                    1470 non-null   int64
28  TotalWorkingYears                   1470 non-null   int64
29  TrainingTimesLastYear               1470 non-null   int64
30  WorkLifeBalance                     1470 non-null   int64
31  YearsAtCompany                      1470 non-null   int64
32  YearsInCurrentRole                  1470 non-null   int64
33  YearsSinceLastPromotion              1470 non-null   int64
34  YearsWithCurrManager                 1470 non-null   int64
dtypes: int64(26), object(9)
memory usage: 402.1+ KB

```

heatmap to check the missing value

```
In [10]: plt.figure(figsize =(10, 4))  
sns.heatmap(dataset.isnull(), yticklabels = False, cbar = False, cmap ='viridis')
```


```
Out[10]: <matplotlib.axes._subplots.AxesSubplot at 0x27adf548e80>
```



So, we can see that there are no missing values in the dataset.


```
In [11]: sns.set_style('darkgrid')  
sns.countplot(x = 'Attrition', data = dataset)
```

```
Out[11]: <matplotlib.axes._subplots.AxesSubplot at 0x27adf5f98b0>
```




```
In [12]: sns.lmplot(x = 'Age', y = 'DailyRate', hue = 'Attrition', data = dataset)
```

```
Out[12]: <seaborn.axisgrid.FacetGrid at 0x27adf5d68b0>
```



```
In [13]: plt.figure(figsize =(10, 6))  
sns.boxplot(y ='MonthlyIncome', x ='Attrition', data = dataset)
```

```
Out[13]: <matplotlib.axes._subplots.AxesSubplot at 0x27adff9c640>
```



In the dataset there are 4 irrelevant columns, i.e:EmployeeCount, EmployeeNumber, Over18 and StandardHour. So, we have to remove these for more accuracy.

```
In [21]: dataset.drop('EmployeeCount', axis = 1, inplace = True)  
dataset.drop('StandardHours', axis = 1, inplace = True)  
dataset.drop('EmployeeNumber', axis = 1, inplace = True)  
dataset.drop('Over18', axis = 1, inplace = True)  
print(dataset.shape)
```

```
(1470, 31)
```

```
In [22]: y = dataset.iloc[:, 1]  
X = dataset  
X.drop('Attrition', axis = 1, inplace = True)
```

```
In [23]: from sklearn.preprocessing import LabelEncoder
lb = LabelEncoder()
y = lb.fit_transform(y)
```

In the dataset there are 7 categorical data, so we have to change them to int data, i.e we have to create 7 dummy variable for better accuracy.

```
In [25]: dum_BusinessTravel = pd.get_dummies(dataset['BusinessTravel'],prefix = 'BusinessTravel')
dum_Department = pd.get_dummies(dataset['Department'],prefix = 'Department')
dum_EducationField = pd.get_dummies(dataset['EducationField'],prefix = 'EducationField')
dum_Gender = pd.get_dummies(dataset['Gender'], prefix = 'Gender', drop_first = True)
dum_JobRole = pd.get_dummies(dataset['JobRole'],prefix = 'JobRole')
dum_MaritalStatus = pd.get_dummies(dataset['MaritalStatus'],prefix = 'MaritalStatus')
dum_OverTime = pd.get_dummies(dataset['OverTime'], prefix = 'OverTime', drop_first = True)
# Adding these dummy variable to input X
X = pd.concat([X, dum_BusinessTravel, dum_Department,dum_EducationField, dum_Gender, dum_JobRole,dum_MaritalS
tatus, dum_OverTime], axis = 1)
# Removing the categorical data
X.drop(['BusinessTravel', 'Department', 'EducationField', 'Gender', 'JobRole', 'MaritalStatus', 'OverTime'],
axis = 1, inplace = True)

print(X.shape)
print(y.shape)

(1470, 49)
(1470,)
```

```
In [26]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.25, random_state = 40)
```

Apply KNN Algo



```
In [27]: from sklearn.neighbors import KNeighborsClassifier
neighbors = []
cv_scores = []

from sklearn.model_selection import cross_val_score
# perform 10 fold cross validation
for k in range(1, 40, 2):
    neighbors.append(k)
    knn = KNeighborsClassifier(n_neighbors = k)
    scores = cross_val_score(knn, X_train, y_train, cv = 10, scoring = 'accuracy')
    cv_scores.append(scores.mean())
error_rate = [1-x for x in cv_scores]

# determining the best k
optimal_k = neighbors[error_rate.index(min(error_rate))]
print('The optimal number of neighbors is % d ' % optimal_k)

# plot misclassification error versus k
plt.figure(figsize = (10, 6))
plt.plot(range(1, 40, 2), error_rate, color='blue', linestyle='dashed', marker='o', markerfacecolor='red',
markersize = 10)
plt.xlabel('Number of neighbors')
plt.ylabel('Misclassification Error')
plt.show()
```

The optimal number of neighbors is 7



The optimal number of neighbors is 7

```
In [28]: from sklearn.model_selection import cross_val_predict, cross_val_score
from sklearn.metrics import accuracy_score, classification_report
from sklearn.metrics import confusion_matrix

def print_score(clf, X_train, y_train, X_test, y_test, train = True):
    if train:
        print("Train Result:")
        print("-----")
        print("Classification Report: \n {}".format(classification_report(
            y_train, clf.predict(X_train))))
        print("Confusion Matrix: \n {}".format(confusion_matrix(
            y_train, clf.predict(X_train))))

        res = cross_val_score(clf, X_train, y_train, cv = 10, scoring = 'accuracy')
        print("Average Accuracy: \t {:.4f}".format(np.mean(res)))
        print("Accuracy SD: \t\t {:.4f}".format(np.std(res)))
        print("accuracy score: {:.4f}\n".format(accuracy_score(y_train, clf.predict(X_train))))
        print("-----")

    elif train == False:
        print("Test Result:")
        print("-----")
        print("Classification Report: \n {}".format(classification_report(y_test, clf.predict(X_test))))
        print("Confusion Matrix: \n {}".format(confusion_matrix(y_test, clf.predict(X_test))))
        print("accuracy score: {:.4f}\n".format(accuracy_score(y_test, clf.predict(X_test))))
        print("-----")

knn = KNeighborsClassifier(n_neighbors = 7)
knn.fit(X_train, y_train)
print_score(knn, X_train, y_train, X_test, y_test, train = True)
print_score(knn, X_train, y_train, X_test, y_test, train = False)
```

Train Result:

Classification Report:

	precision	recall	f1-score	support
0	0.86	0.99	0.92	922
1	0.83	0.19	0.32	180
accuracy			0.86	1102
macro avg	0.85	0.59	0.62	1102
weighted avg	0.86	0.86	0.82	1102

Confusion Matrix:

```
[[915  7]
 [145 35]]
```

Average Accuracy: 0.8421

Accuracy SD: 0.0148

accuracy score: 0.8621

Test Result:

Classification Report:

	precision	recall	f1-score	support
0	0.84	0.96	0.90	311
1	0.14	0.04	0.06	57
accuracy			0.82	368
macro avg	0.49	0.50	0.48	368
weighted avg	0.74	0.82	0.77	368

Confusion Matrix:

```
[[299 12]
 [ 55  2]]
```

accuracy score: 0.8179

In []: